

SURANJIT BANIK

Data & Intelligence Engineer | Azure Databricks Specialist

Toronto, ON | 647 616 9494 | suran.cse@gmail.com
linkedin.com/in/suranjitbanik | mastermindlab.github.io

PROFESSIONAL SUMMARY

Data & Intelligence Engineer with 5+ years building on Azure Databricks: Delta Lake, Lakeflow Declarative Pipelines, Unity Catalog, Databricks Apps, Asset Bundles, and Genie. Delivers production lakehouse pipelines, AI powered data products, and CI/CD automation across insurance, digital media, and analytics. Early adopter of agentic development on Databricks using Claude Code, the Databricks MCP server, and Omnigent.

CORE SKILLS

Databricks: Delta Lake, Lakeflow Declarative Pipelines (DLT), Unity Catalog, Databricks Apps, Workflows, Asset Bundles, Auto Loader, Genie, Lakebase

Languages: Python, PySpark, Spark SQL, SQL

Cloud: Azure (ADLS, DevOps, ADF, Logic Apps), AWS S3, GCP BigQuery

AI / Agentic Tooling: Claude Code, Databricks MCP server, Databricks CLI, Omnigent, Claude Cowork

Tools: GitHub, Jira, Confluence, Streamlit

Strengths: Medallion Architecture, Metadata Driven Ingestion, Data Governance, CI/CD, ETL Design, Performance Tuning

PROFESSIONAL EXPERIENCE

Data & Intelligence Engineer | Stratis

August 24, 2026 - Present

- Building the agentic AI platform for MarTech and large banks.

Senior Data Engineer | Zurich Canada, Toronto, ON

April 2024 - August 2026

- Built an AI powered policy and claims app with Databricks Genie and Lakebase, letting underwriters query live data in natural language and persist open claims notes to a Lakebase OLTP store, cutting insight turnaround from days to seconds.
- Developed a Databricks App for bordereaux (BDX) file intake with Claude LLM powered column mapping that strips special characters, converts headers to snake_case, and auto maps columns to the target schema, eliminating manual mapping fixes in lakehouse ingestion.
- Built declarative medallion pipelines (bronze, silver, gold) on Lakeflow with data quality expectations, CDC, Unity Catalog governance, and Auto Loader for schema resilient ingestion across insurance data platforms.

SURANJIT BANIK

PROFESSIONAL EXPERIENCE / CONTINUED

Senior Data Engineer | Zurich Canada, Toronto, ON

April 2024 - August 2026 (continued)

- Designed config driven PySpark ETL frameworks reused across 30+ brokers data, cutting new source onboarding time by roughly 60%.
- Modernized CI/CD with Databricks Asset Bundles and Azure DevOps across dev, qa, uat, and prod, improving release velocity and environment portability.
- Optimized Spark jobs with Predictive Optimization, liquid clustering (CLUSTER BY) on serverless compute, and join strategy tuning, reducing runtimes by roughly 40%.
- Adopted agentic engineering workflows pairing Claude Code with the Databricks MCP server and CLI to scaffold pipelines, generate tests, and debug jobs, shipping multiple projects ahead of schedule.

Data Engineer | GroupM, Toronto, ON

August 2021 - April 2024

- Built Databricks and PySpark pipelines on GCP integrating BigQuery, Bigtable, and Cloud Functions, improving reporting performance by 60%.
- Delivered multi cloud solutions across GCP, Azure, and AWS and led data governance and compliance for enterprise data warehousing.

Data Engineer | Clue, Toronto, ON

January 2021 - July 2021

- Designed graph data models with Neo4j and Cypher and built ETL pipelines on AWS (S3, Athena, SageMaker, DataBrew) for digital marketing analytics.

Data Analyst | Clue, Toronto, ON

February 2020 - December 2020

- Built Azure analytics services ingesting Salesforce data for media campaign reporting with near zero downtime on real time pipelines.

Digital Analyst | Mindshare, Toronto, ON

May 2019 - February 2020

- Architected Datorama solutions for Nestle and CPG clients, streamlining marketing data integration and campaign ROI measurement.

EDUCATION

Master of Engineering, Computer Engineering | University of Ottawa

Bachelor of Engineering, Computer Science | Visvesvaraya Technological University

CERTIFICATIONS & PROFESSIONAL DEVELOPMENT

Databricks Certified Data Engineer Associate

AWS Certified Cloud Practitioner | Valid through December 2026

Databricks Data + AI Summit, San Francisco | 2026

Hands-on training on Lakeflow, Lakebase, and Genie